

Modeling AIE-ML for Matrix Multiplication

Elliott Binder, Advisor: Tze Meng Low
Electrical and Computer Engineering

Carnegie
Mellon
University

Introduction

Analytically modeling **AIE-ML** architectures to **quickly and efficiently** map applications to compute and memory resources.

AIE-ML is a spatial architecture that can be found in mobile and desktop AMD processors as XDNA NPUs.

AIE-ML **differs** from Versal AIE architecture:

- Low precision (e.g. int16, int8) matrix instructions
- No programmable logic fabric
- Memory tile with more capacity and connectivity

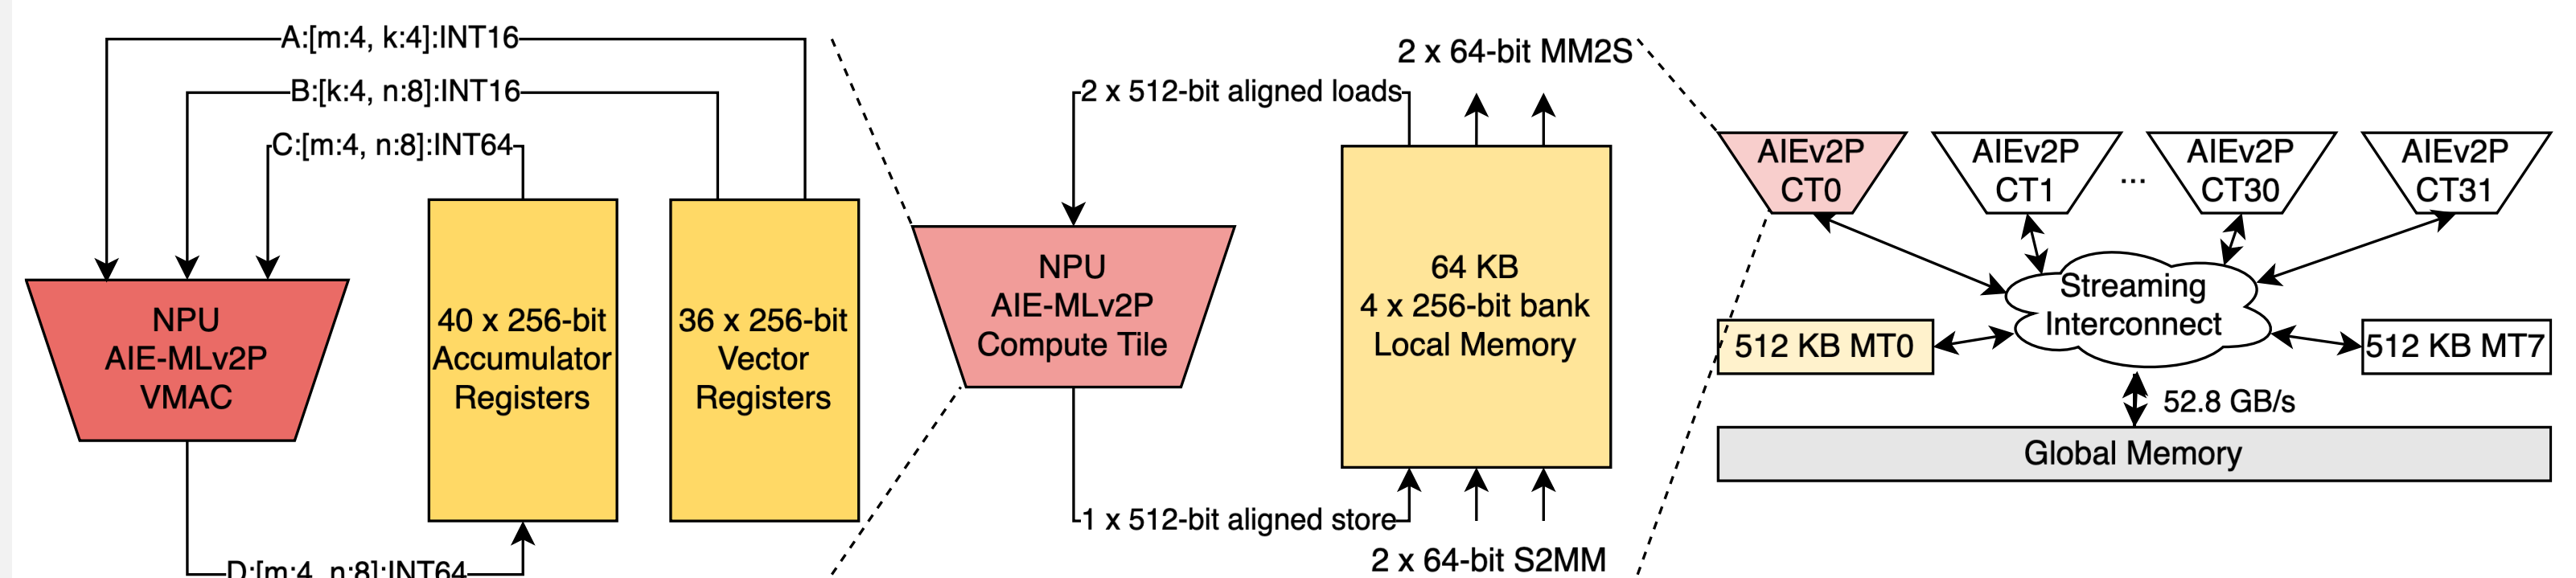
XDNA 2 NPUs **improve** over previous generation:

- Higher compute throughput
- Higher interconnect bandwidth
- More compute and memory tiles


Methods

Leverage models of **hardware resources**:

- Functional unit throughput and latency
- Register capacity and layouts for MMAC instructions
- Local memory structure, throughput and latency
- Bandwidth, connectivity of Compute/Memory tiles
- Bandwidth available from main memory to simplify development effort of AIE-ML applications.



With Support From


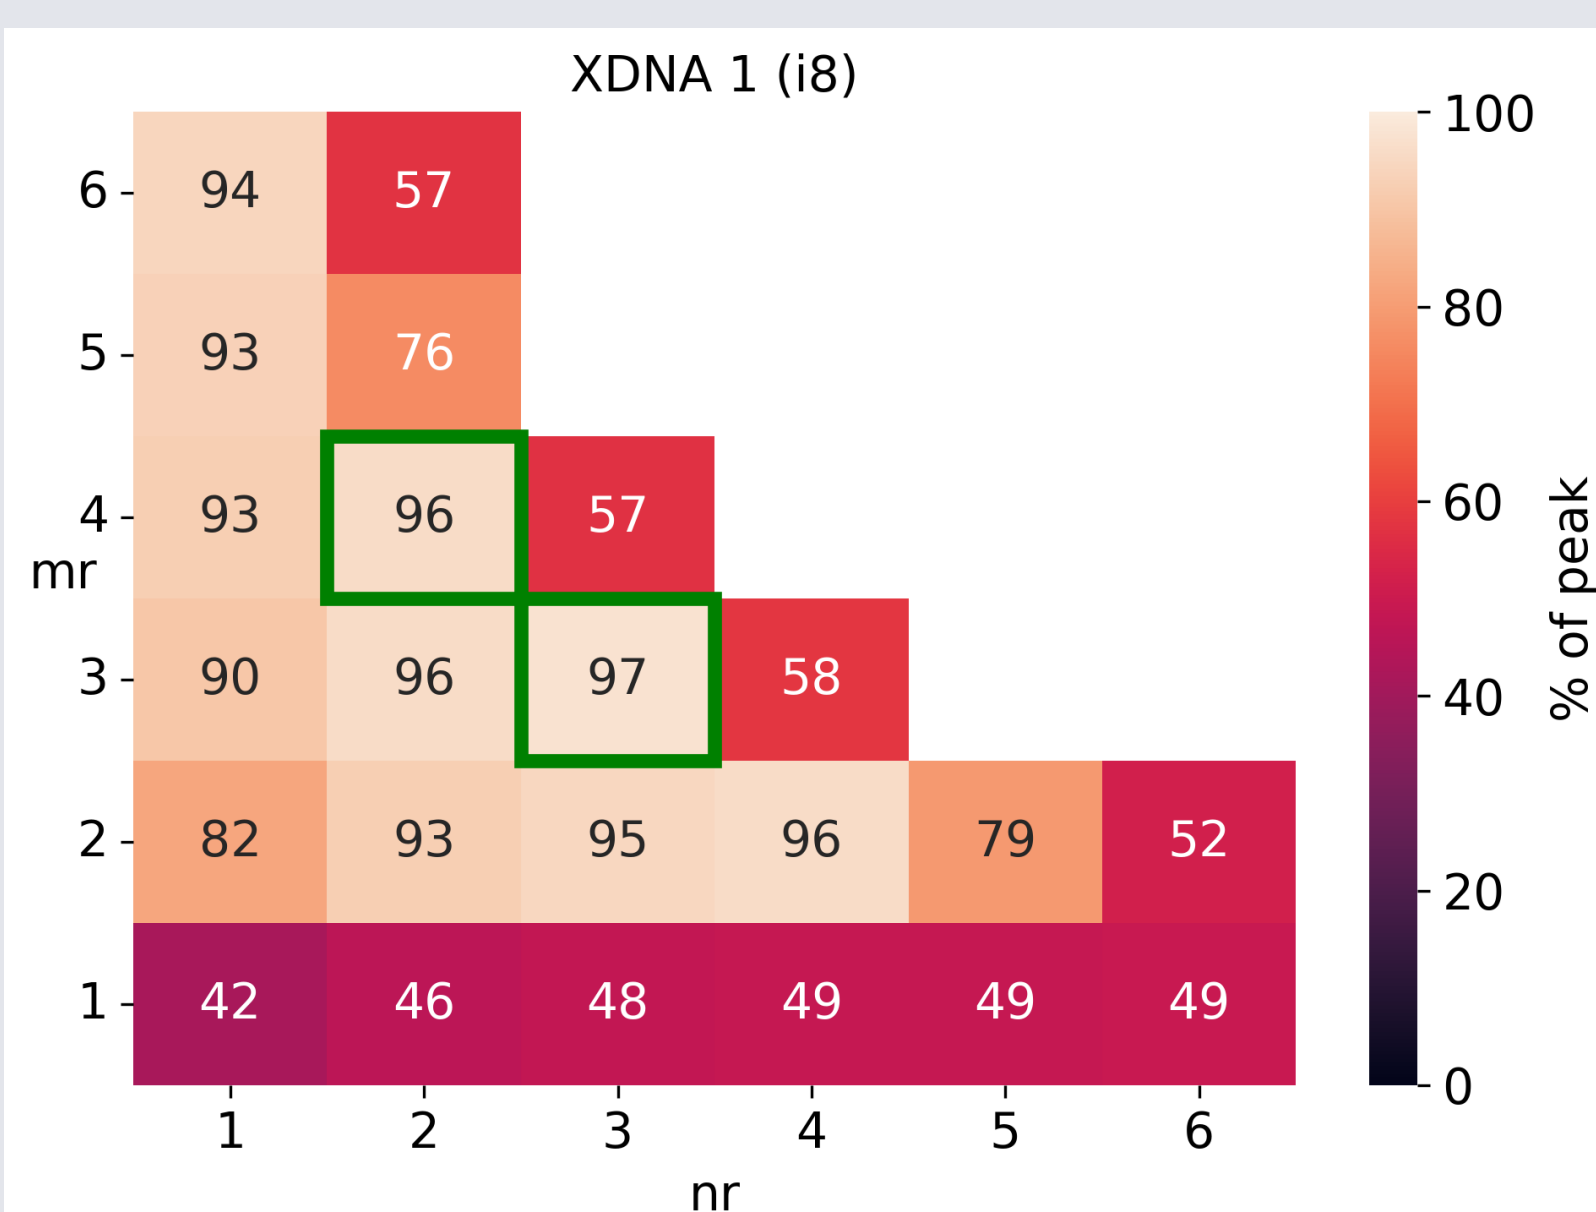


Results

Modeling the Register Tile Size

Register tile should

- Hide latency of MMAC
- Hide loads from local memory
- Fit inputs into vector registers
- Fit output into acc. registers



Parameters satisfying models achieve 94-97% of peak compute throughput.


Modeling the Compute Tile Size

Local memory tile should

- hide cost of streaming interface
- fit within local memory capacity

Insights:

- Input stationary enables more efficient kernel via larger K tile factor
- Single-buffer stationary matrix to free capacity for larger, more efficient tiles



Choosing tile sizes just large enough to hide memory cost can enable larger K tiling values which yield higher kernel efficiency.

Poster



Contact



Modeling the Array Tile Size

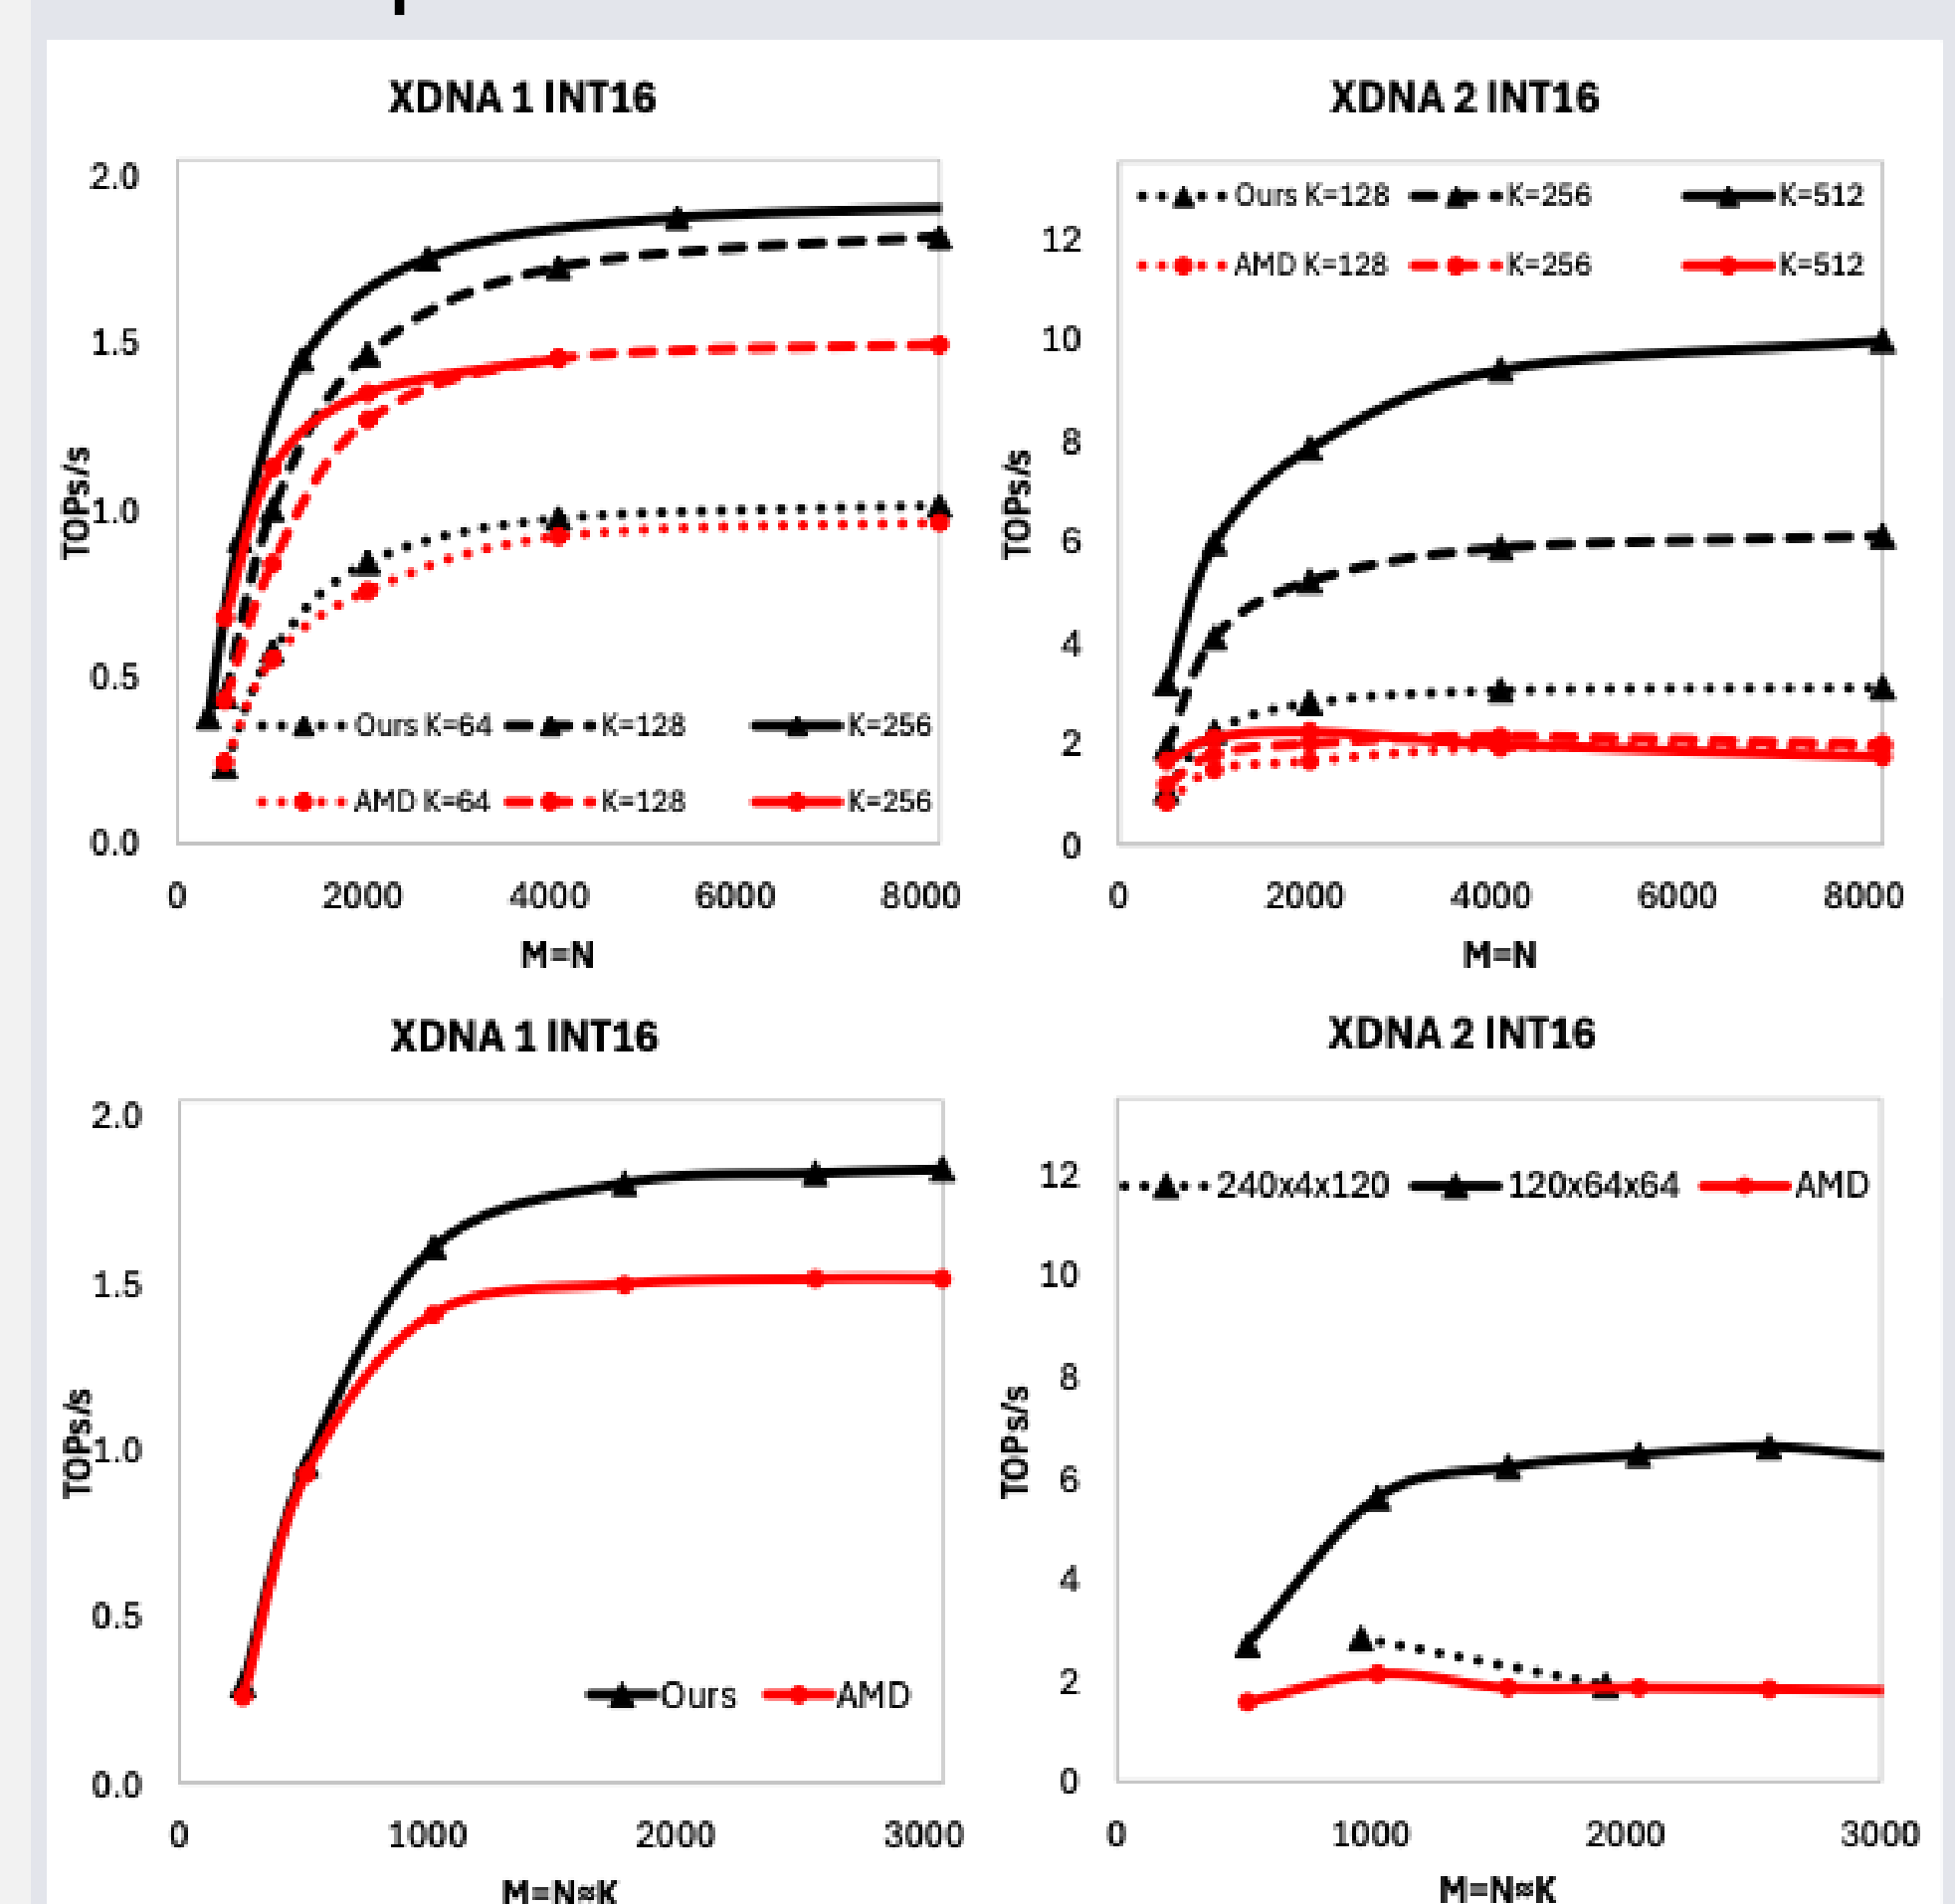
Matrices are staged in memory tiles between main memory and the array.

The array tile size should

- hide cost of accessing non-resident matrices in main memory

Insights:

- Input stationary algorithm more effectively uses R+W bandwidth
- Opposing skew of local tile size and parallelization should be used



Considering main memory bandwidth yields higher performance than SOTA.

Discussion

In some scenarios, our models lead to implementations that can achieve 90+% of compute peak even when matrices reside in main memory. Yet, other scenarios present no theoretically viable solution capable of achieving this performance.

Several improvements we'd like to see for the AIE-ML programming interface include:

- Convenient way of passing more than two arguments to a compute tile, currently limited by static mapping to streams
- Outer-loop reuse of data in memory tiles for more effective tiling at the array-level, to alleviate limited main mem. BW
- Manual control over VLIW instruction generation, including scheduling and register allocation, to improve kernel
- Larger multi-dimensional data transfers from main memory to streaming interconnect, currently limiting problem sizes